

Laboratory 1

Yrkkö Äkkijyrkkä yrkko@kth.se Martin Yrjölä yrjola@kth.se

Introduction

We have been collaborating with a group of other students to automate the Simics benchmarking in the git repository published at https://github.com/myrjola/comparch_labs. The data used for our figures and tables and the code to generate them are published in the same repository. The majority of the lab time has been spent developing the automation and doing data gathering. We have discussed our findings actively with the repository collaborators, but also took care to conduct our own benchmarks and write our own analysis.

3.2 Collecting statistics from simple cache

The performance results are given in tbl. 1. It is clear that the parser benchmark has the best cache performance because it has the highest hit rate. The vortex benchmark has the worst performance only scoring a hit 56.9 percent of the time.

Table 1: Cache performance with default settings

Benchmark	Inst hit rate	Read hit rate	Write hit rate	Hit rate
parser	92.8%	87.2%	63.4%	84.3%
equake	79.3%	87.1%	80.9%	81.6%
vortex	58.7%	74.8%	28.2%	56.9%

3.3 Determining benchmark working set size

By increasing the cache size the amount of conflict and capacity misses will decrease. Because the cache has been warmed up in our benchmark, there should be very few compulsory misses. The cache size at the point when the hit rate stabilizes should give us an approximation of the program's working set. A




Figure 1: Hit rate for different cache sizes and degrees of associativity

visualization of the hit rate is presented in fig. 1. We also tested the performance by increasing the associativity to 4 with the aim of decreasing conflict misses in order to get a better approximation. The graph suggests that the parser benchmark has the smallest working set at approximately 60kB. The working set of the vortex benchmark is around 2MB. The equake benchmark has the largest working set. We estimate it is approximately 4MB.

3.4 Minimal instruction and data caches

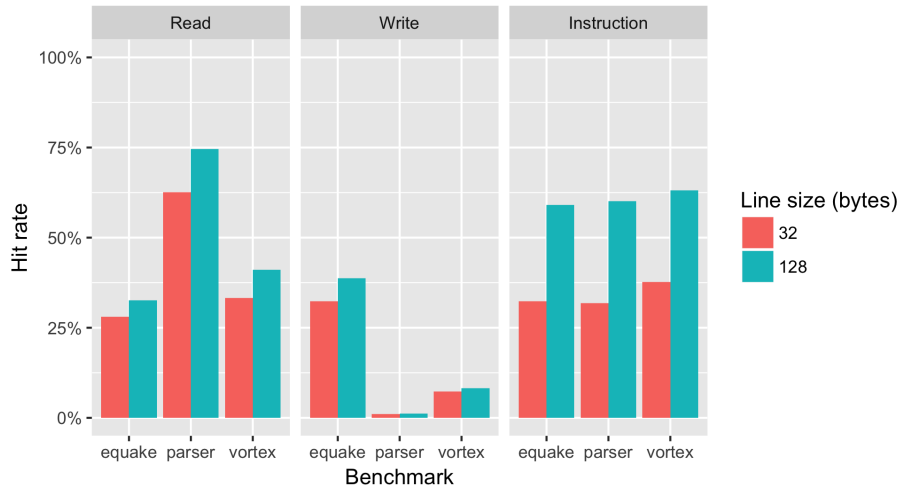


Figure 2: Benchmark performance with one cache line

The performance using one cache line with different sizes is presented in fig. 2. Using one cache line will give us a hint of the spatial locality of the data. A line size of 32 bytes is quite small and will only perform well if the data is accessed sequentially. My intuition says that code is accessed mostly in this manner. The program counter will keep on increasing until a branch happens. The UltraSparc T1 processor has instructions of 32 bits, so without branches we should see a hit rate of 7/8 with a line size of 32 bits and a 31/32 hit rate for a line size of 128 (*UltraSPARC Architecture 2005*, n.d.). However, the performance doesn't come close to these values. Maybe there's more branches in the benchmarks than in typical code. One thing to note is that the instruction fetch hit rate almost doubles with the larger line size.

The data reads and writes share the same cache, so it is very likely that they will interfere with each other. We see a big variance between the benchmarks. The parser has the most sequential reads, but the worst write performance. The equake benchmarks has the best write performance.

The question was about comparing relative locality of data against instruction accesses. The data suggests that instruction accesses tend to follow the principles of temporal and spatial locality quite well. Data access is very much dependent on the benchmark. The parser benchmark achieves higher read hit rate than the code, but both the parser and vortex suffer of bad write performance.

3.5 Collecting statistics from a cache hierarchy

The performance for the cache hierarchy for different sizes of L2 cache is given in fig. 3. The first question was to assess how effective the L2 cache is at collecting misses from the L1 level. We would argue that the L2 cache performs very well, because the hit rate is very high for every benchmark.

The second question was related to how the size of the L2 cache changes its performance. We choose to analyze the performance in the vortex benchmark, which has the most variance in the L2 results. The small size of 8 KiB has significantly lower hit rate than the 500 kiB and 1 MiB L2 caches. The difference between 500 kiB and 1 MiB is negligible, so it doesn't make sense to invest in a larger L2 cache for this benchmark. Choosing between 8 kiB and 500 kiB is a tradeoff between cost and performance. In the next section we will give a useful metric for assessing the performance between the different sizes.

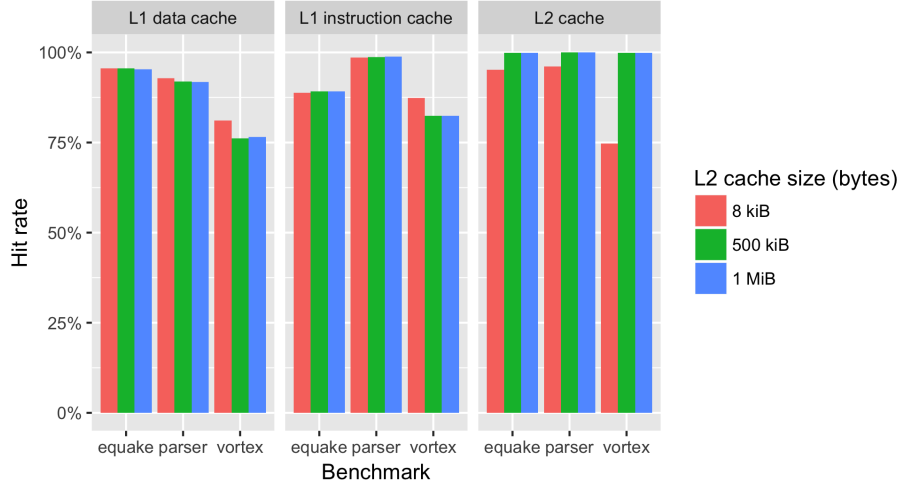


Figure 3: Cache hierarchy performance

Next we will calculate the average memory access time. We will use the following formula to calculate it.

$$3 + L1_{\text{miss rate}} * (10 + L2_{\text{miss rate}} * 200)$$

The average access times are given in tbl. 2. We calculated them by aggregating the statistics for all the benchmarks. As expected the access time gets shorter by increasing the cache size.

Table 2: Average memory access time for the cache hierarchies

L2 cache size	Average memory access time (cycles)
8 kiB	6.1996
500 kiB	4.1795
1 MiB	4.1769

3.6 Application tuning for a given cache hierarchy

Table 3: Average memory access time for the gemm benchmark

AMAT (cycles)	L2 lines	DC line size	DC lines	IC lines	DC assoc	L2 assoc	DC hit rate	IC hit rate	L2 hit rate
10.69258	32	16	16	8	1	4	0.88	0.69	0.84
10.69258	32	16	16	8	1	4	0.88	0.69	0.84
11.02387	32	64	16	8	1	4	0.85	0.69	0.85
11.02387	32	64	16	8	1	4	0.85	0.69	0.85
11.13431	32	32	16	8	1	4	0.84	0.69	0.85
11.13431	32	32	16	8	1	4	0.84	0.69	0.85
11.26225	16	16	16	8	1	4	0.88	0.69	0.83
11.26225	16	16	16	8	1	4	0.88	0.69	0.83
11.37848	32	64	8	8	1	4	0.82	0.69	0.86
11.37848	32	64	8	8	1	4	0.82	0.69	0.86
11.53585	32	32	8	8	1	4	0.81	0.69	0.86
11.53585	32	32	8	8	1	4	0.81	0.69	0.86
11.56691	16	64	16	8	1	4	0.85	0.68	0.84
11.56691	16	64	16	8	1	4	0.85	0.68	0.84
11.73972	16	32	16	8	1	4	0.84	0.68	0.84
11.73972	16	32	16	8	1	4	0.84	0.68	0.84
11.90727	32	64	4	8	1	4	0.78	0.69	0.87
11.90727	32	64	4	8	1	4	0.78	0.69	0.87
12.04599	16	64	8	8	1	4	0.82	0.68	0.84
12.04599	16	64	8	8	1	4	0.82	0.68	0.84
12.18663	16	32	8	8	1	4	0.81	0.68	0.85
12.18663	16	32	8	8	1	4	0.81	0.68	0.85
12.29470	8	16	16	8	1	4	0.88	0.68	0.80
12.29470	8	16	16	8	1	4	0.88	0.68	0.80

AMAT (cycles)	L2 lines	DC line size	DC lines	IC lines	DC assoc	L2 assoc	DC hit rate	IC hit rate	L2 hit rate
12.43491	8	64	16	8	1	4	0.84	0.68	0.82

In choosing the most suitable cache we have to take both performance and cost metrics into account. Most of the performance metrics can be generated by Simics, but we plan to use Cacti 5.3 to gain access time, power and area overheads. We use the default settings from the Cacti web interface. What we found out by testing Cacti is that it doesn't work for small caches. At least the results seem to be illogical with access times of 1e39 ns for many of our cache configurations. Because of this bug we were unable find the optimal configuration with regard to all metrics.

We performed 584 tests and chose the average memory access time as the most suitable performance metric. We calculate the AMAT in the same manner as in 3.5 with the same miss penalties, we would have liked to use dynamic miss penalties based on the cache configuration, but that was not possible due to the Cacti bugs discussed earlier. The gemm benchmark has been run 100 million cycles after the first magic breakpoint.

The top-25 best performing configurations are given in tbl. 3. The blocked matrix multiply works well with very small caches. Even the L2 cache can be very small without too much performance degradation. Without being able to assess all metrics with Cacti we have to hazard a guess on the most balanced configuration. We would like to choose the 23rd best configuration for the application, because that is the best one with an L2 cache of only 8 lines and still with good performance.

4.4 Multithreaded performance study

The miss rates for the multithreaded benchmark are given in fig. 4. The results for the L1 cache are averaged over all the caches. What we can observe is that the miss rates decrease substantially for the L1 cache when we quadruple its size, this is expected because there is likely a lot of conflict misses with the smaller cache. The miss rates for L2 instruction fetch and write increases when L1 cache is enlarged. The main reason is that earlier the L1 conflict misses got served by the L2 cache but with the larger L1 cache most of the conflict misses become hits and the thus there are less hits overall in the L2 cache.

MESI statistics for the L1 caches are given in fig. 5. The values are the arithmetic mean of all the statistics on the different processors. The L1 cache is write through, therefore there are no 'modified'-transactions. The exclusive to shared statistics are increased because it is more likely that any given address is available

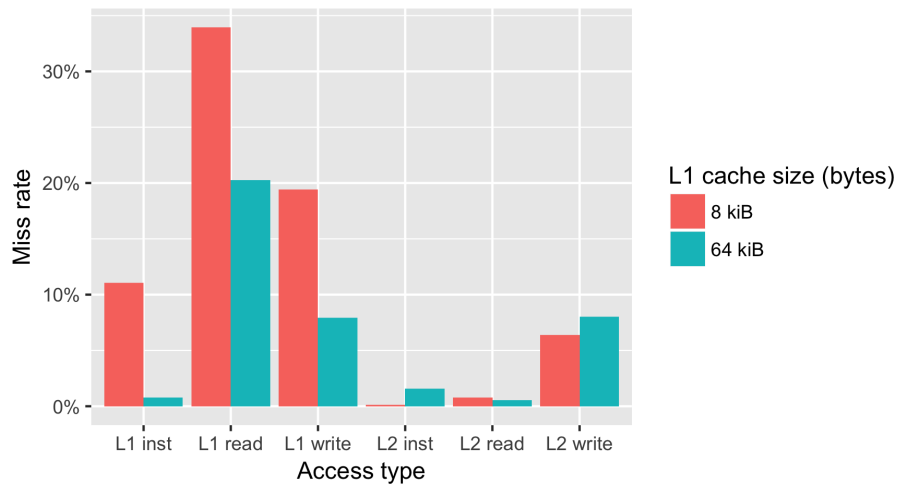


Figure 4: Multithreaded cache performance

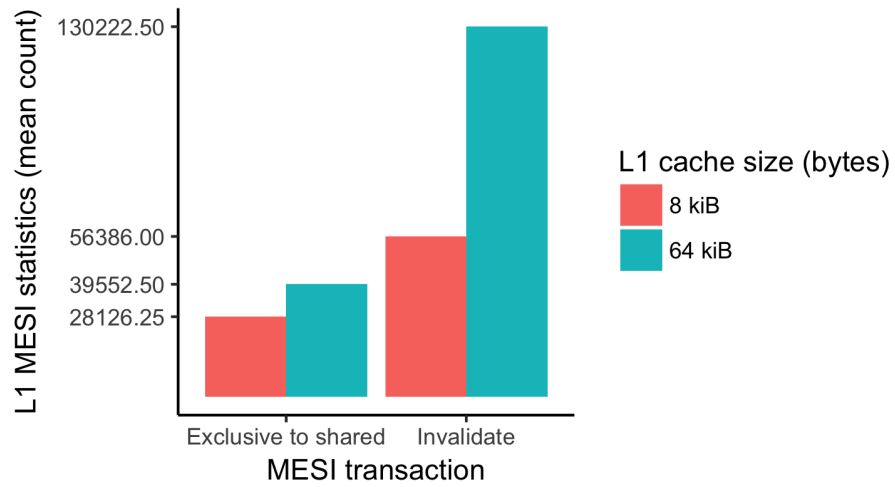


Figure 5: Multithreaded MESI statistics

in the larger caches. This also leads to a dramatic increase in invalidates because more data is shared in a larger cache and thus a write has a higher probability to invalidate more data than with the smaller cache.

5.5 Critical sections

1. The data is most of the time inconsistent when using `./lab1.5 -t critical -c pthreads`, increasingly so when adding threads to the execution. This is explained by race conditions because there is no synchronization. When one thread has read the counter value it should be locked from reading and writing by other threads until the value has been incremented/decremented.
2. Adding the `enter_critical()` and `exit_critical()` makes the data consistent across executions by using mutexes.
3. Dekker's algorithm
 - a. The variables are marked volatile in order to protect false writes from concurrently accessing threads because of compiler optimizations that would not create problems in a single-threaded execution.
 - b. The CPU's out-of-order execution causes the inconsistency of data.
4. Memory fence Implementation, in `enter_critical()`:

```
if(thread == 1){
    flag[1] = true;
    MFENCE();
    while(flag[0]){
        if(turn != 1){
            flag[1] = false;
            while(turn != 1){
                //Busy wait
            }
            flag[1] = true;
        }
    }
}
MFENCE();
```

5. Atomic increment & decrement

Using the non-atomic instructions, the data is inconsistent. This is because there is no protective methods for the data. The memory location can be accessed by another thread while it is in the middle of the other CPU's pipeline, containing the old value before the other CPU has written into yet.

Using atomic instructions, specifically the `lock`-instruction, a data atomicity is ensured by system-level methods. On older CPUs of the Intel family, the instruction locked the memory bus for the duration of the operation, but from P6 forward, the instruction enforces cache-locking. This guarantees exclusive ownership of the cache-line for this data during the operation. Then the cache-coherency mechanisms will ensure that the data is correct after the operation. (*Intel® 64 and Ia-32 Architectures Software Developer’s Manual*, n.d.)

6. Atomic compare & exchange

Using non-atomic instructions, the data is yet again inconsistent. Mostly for same reasons as described above. Using CAS here is also the risk of an ABA-problem, where the other thread can write memory twice so that the resulting bits are same but the counter value does not keep track of this. Using the `lock`-version of `cmpxchgl` results in consistent data. It ensures that the data being read for comparison has not changed meanwhile.

7. Runtime comparison, critical vs. atomic operations

As the data in tbl. 4 shows, critical sections are much slower. The time for Pthreads would vary significantly over different execution rounds, the time presented in the table is worst case recorded. We can see that instruction-level methods are much faster than software-level protective methods. They both induce stalling in the CPUs but instruction-level methods can reduce the stalls to be only the length of the instruction.

8. Runtime comparison, atomic vs. non-atomic operations

Comparing the instruction-level methods in tbl. 4, we can see that atomic instructions are slower as they induce stalling, but are a necessary for data consistency. CAS operations are slower but they offer more flexibility as they are not limited to increment/decrement operations, and could be used to implement complicated functions with good performance and data integrity.

Table 4: Runtime performance comparison table

Property	Pthreads	Dekker	Atomic inc/dec	Non- atomic inc/dec	Atomic CAS	Non- atomic CAS
Average execution time:	1.2349 s	0.3516 s	0.0430 s	0.0100 s	0.1139 s	0.0332 s

Bibliography

Intel® 64 and Ia-32 Architectures Software Developer's Manual. n.d. http://www.intel.com/Assets/en_US/PDF/manual/253668.pdf: Intel Corporation.

UltraSPARC Architecture 2005. n.d. <http://www.oracle.com/technetwork/systems/opensparc/t1-06-ua2005-d0-9-2-p-ext-1537734.html>: Sun microsystems.